

## Data Mining Assignment 4

1) Read Chapter 4 (all sections) and Chapter 5 (Sections 5.2, 5.5, 5.6 and 5.7).

2) Consider the following data set for a binary class problem.

A	B	Class Label
T	F	+
T	T	+
T	T	+
T	F	-
T	T	+
F	F	-
F	F	-
F	F	-
T	T	-
T	F	-

Calculate the misclassification error rate when splitting on A and B to determine the best split. Which of these splits considered is the best according to misclassification error rate?

**Splitting on A:**

	A = T	A = F
+	4	0
-	3	3

Misclassification error of A:

$$\text{Rate} = 1 - [(\text{TP} + \text{TN}) / \text{Total records}]$$

$$= 1 - [(4 + 3) / 10]$$

$$= 3/10 = \mathbf{0.3}$$

**Splitting on B:**

	B = T	B = F
+	3	1
-	1	5

Misclassification error of B:

$$\text{Rate} = 1 - [(\text{TP} + \text{TN}) / \text{Total records}]$$

$$= 1 - [(3 + 5) / 10]$$

$$= 2/10 = \mathbf{0.2}$$

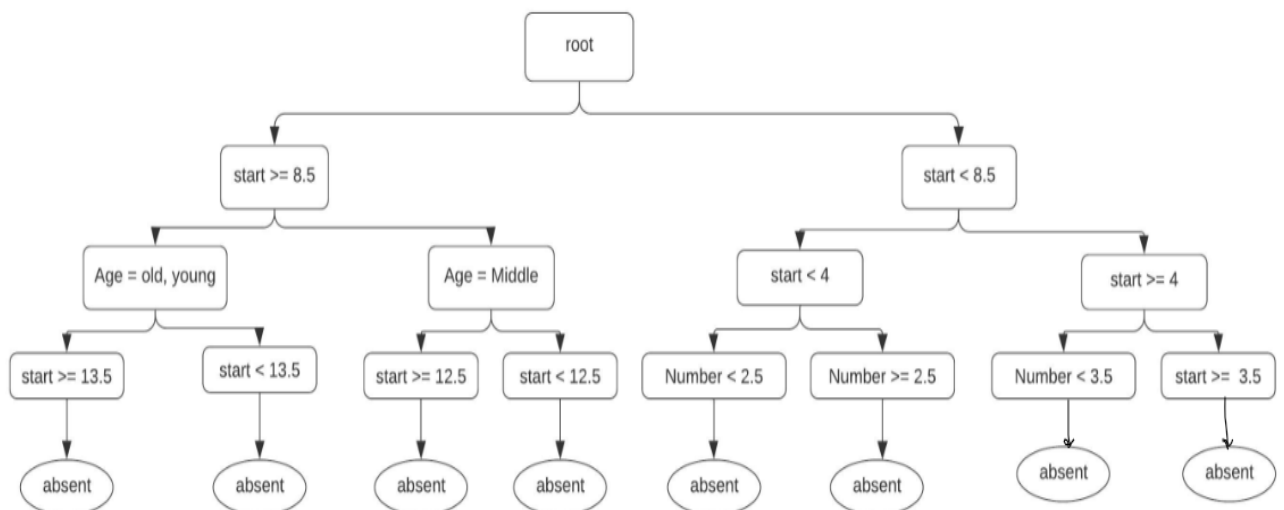
**As the misclassification error of B is lower, hence the splitting the table on B should be considered as the best split.**

3) Consider the training examples shown below for a binary classification problem.

Instance	$a_1$	$a_2$	$a_3$	Target Class
1	T	T	1.0	+
2	T	T	6.0	+
3	T	F	5.0	-
4	F	F	4.0	+
5	F	T	7.0	-
6	F	T	3.0	-
7	F	F	8.0	-
8	T	F	7.0	+
9	F	T	5.0	-

For  $a_3$ , which is a continuous attribute compute misclassification error rate for every possible split to determine the best split. Which of these splits considered is the best according to misclassification error rate?

4) The file [http://www-stat.wharton.upenn.edu/~dmease/rpart\\_text\\_example.txt](http://www-stat.wharton.upenn.edu/~dmease/rpart_text_example.txt) gives an example of text output for a tree fit using the `rpart()` function in R from the library `rpart`. Use this tree to predict the class labels for the 10 observations in the test data [http://www-stat.wharton.upenn.edu/~dmease/test\\_data.csv](http://www-stat.wharton.upenn.edu/~dmease/test_data.csv) linked here. Do this manually - do not use R or any software.



Age	Number	Start	
Middle	5	10	present
Young	2	17	absent
Old	10	6	present
Old	4	15	absent
Middle	5	15	absent
Young	3	13	absent
Old	5	8	present
Young	7	8	absent
Middle	3	13	absent

5) I split the popular sonar data set into a training set ([http://www-stat.wharton.upenn.edu/~dmease/sonar\\_train.csv](http://www-stat.wharton.upenn.edu/~dmease/sonar_train.csv)) and a test set ([http://www-stat.wharton.upenn.edu/~dmease/sonar\\_test.csv](http://www-stat.wharton.upenn.edu/~dmease/sonar_test.csv)). Use R to compute the misclassification error rate on the test set when training on the training set for a tree of depth 5 using all the default values except `control=rpart.control(minsplit=0,minbucket=0,cp=-1, maxcompete=0, maxsurrogate=0, usesurrogate=0, xval=0,maxdepth=5)`. Remember that the 61st column is the response and the other 60 columns are the predictors.

```
> sonar_fit <- rpart(y ~ ., x, control = rpart.control(minsplit=0,minbucket=0,cp=-1, ma
xcompete=0, maxsurrogate=0, usesurrogate=0, xval=0,maxdepth=5))
> plot(sonar_fit)
> text(sonar_fit)
> print(sonar_fit)
n= 130
```

```
node), split, n, loss, yval, (yprob)
* denotes terminal node
```

```
1) root 130 64 -1 (0.50769231 0.49230769)
 2) v11>=0.17095 79 21 -1 (0.73417722 0.26582278)
   4) v27>=0.8191 37 2 -1 (0.94594595 0.05405405)
     8) v9>=0.0889 34 0 -1 (1.00000000 0.00000000) *
     9) v9< 0.0889 3 1 1 (0.33333333 0.66666667)
      18) v2< 0.0195 1 0 -1 (1.00000000 0.00000000) *
      19) v2>=0.0195 2 0 1 (0.00000000 1.00000000) *
 5) v27< 0.8191 42 19 -1 (0.54761905 0.45238095)
   10) v54>=0.02075 12 0 -1 (1.00000000 0.00000000) *
   11) v54< 0.02075 30 11 1 (0.36666667 0.63333333)
     22) v8< 0.17045 17 7 -1 (0.58823529 0.41176471)
       44) v36< 0.4491 12 2 -1 (0.83333333 0.16666667) *
       45) v36>=0.4491 5 0 1 (0.00000000 1.00000000) *
     23) v8>=0.17045 13 1 1 (0.07692308 0.92307692)
       46) v1>=0.0925 1 0 -1 (1.00000000 0.00000000) *
       47) v1< 0.0925 12 0 1 (0.00000000 1.00000000) *
 3) v11< 0.17095 51 8 1 (0.15686275 0.84313725)
   6) v52>=0.0209 6 1 -1 (0.83333333 0.16666667)
     12) v1>=0.02225 5 0 -1 (1.00000000 0.00000000) *
     13) v1< 0.02225 1 0 1 (0.00000000 1.00000000) *
   7) v52< 0.0209 45 3 1 (0.06666667 0.93333333)
     14) v19>=0.8351 5 2 -1 (0.60000000 0.40000000)
       28) v26< 0.6153 3 0 -1 (1.00000000 0.00000000) *
       29) v26>=0.6153 2 0 1 (0.00000000 1.00000000) *
     15) v19< 0.8351 40 0 1 (0.00000000 1.00000000) *
```

```
> predictions <- as.numeric(predict(sonar_fit, sonar_test, type = "class"))
> predictions <- replace(predictions, predictions == 1, -1)
> predictions <- replace(predictions, predictions == 2, 1)
> predictions
[1] -1 -1 -1 1 -1 -1 1 1 -1 -1 -1 1 1 1 -1 -1 -1 1 1 -1 -1 1 -1 1 -1
[28] -1 -1 -1 -1 1 -1 1 1 1 1 1 -1 -1 -1 -1 -1 1 -1 1 1 1 -1 -1 1 1
[55] -1 1 1 -1 1 1 1 -1 -1 1 -1 1 -1 1 -1 1 1 -1 -1 1 1 1 1
> actual_values <- sonar_test[, 61]
> actual_values
[1] -1 -1 -1 1 1 1 1 1 -1 -1 -1 1 -1 1 -1 -1 1 1 1 -1 -1 1 -1 -1
[28] -1 1 1 1 1 1 1 1 1 -1 1 -1 -1 -1 1 -1 1 -1 1 1 1 -1 -1 1 -1
[55] -1 -1 -1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 1 -1 -1 1 1 -1 1
```

```
> xtab <- table(predictions, actual_values)
```

```

> results <- confusionMatrix(xtab)
> results
Confusion Matrix and Statistics

              actual_values
predictions -1  1
          -1 33  8
           1 12 25

              Accuracy : 0.7436
              95% CI   : (0.6321, 0.83
    No Information Rate : 0.5769
    P-Value [Acc > NIR] : 0.001668

              Kappa : 0.4831

  Mcnemar's Test P-Value : 0.502335

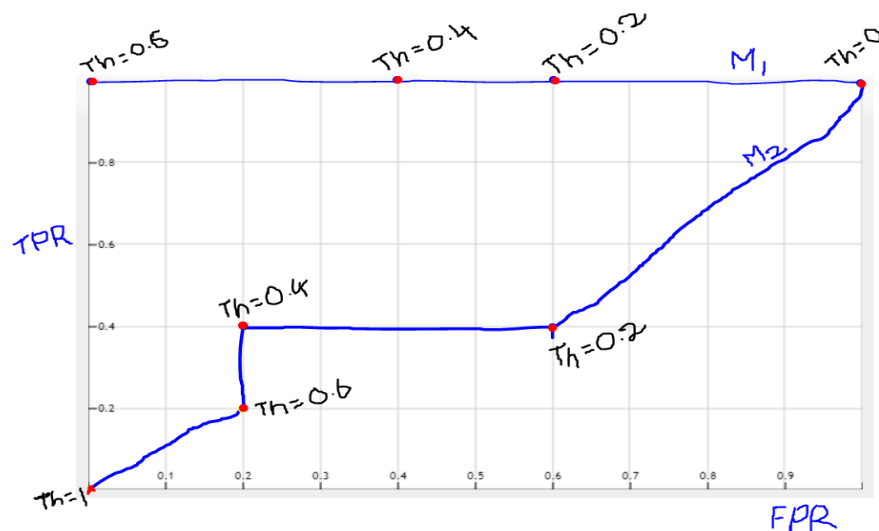
              Sensitivity : 0.7333
              Specificity : 0.7576
              Pos Pred Value : 0.8049
              Neg Pred Value : 0.6757
              Prevalence : 0.5769
              Detection Rate : 0.4231
              Detection Prevalence : 0.5256
              Balanced Accuracy : 0.7455

              'Positive' Class : -1

> mean(predictions != actual_values)
[1] 0.2564103

```

6) Do Chapter 5 textbook problem #17 (parts a and c only) on pages 322-323. Note that there is a typo in part c - it should read "Repeat the analysis for part (b)". We will do part b in class.



### M1

Cut-off threshold = 0.5

TP = 1, 2, 5

TN = 3, 7, 8, 10

FN = 6.9

Precision (p) =  $TP / (TP + FP) \rightarrow 3 / (3+1) = 0.75$

Recall (r) =  $TP / (TP + FN) \rightarrow 3 / (3+2) = 0.6$

The harmonic mean between recall and precision is F1.

$F1 = 2 / (1/r + 1/p) \rightarrow 2 / (1/0.6 + 1/0.75) = 0.9/1.35 = 0.667$

### M2

Cut-off threshold = 0.5

TP = 1

TN = 4, 7, 8, 10

FN = 2, 5, 6, 9

FP = 3

Precision (p) =  $TP / (TP + FP) \rightarrow 1 / (1+1) = 0.5$

Recall (r) =  $TP / (TP + FN) \rightarrow 1 / (1+4) = 0.2$

The harmonic mean between recall and precision is F1.

$F1 = 2 / (1/r + 1/p) \rightarrow 2 / (1/0.2 + 1/0.5) = 0.2/0.7 = 0.2857$

**By comparing M1, M2; F1 is larger for M1 which indicates a better model.**

7) Compute the misclassification error on the training data for the Random Forest classifier to the last column of the sonar training data. Show your R code for doing this.

```
> # Load the party package. It will automatically load other
> # required packages.
> library(party)
> library(randomForest)
>
> # Create the forest.
> output_fit <- randomForest(X.1 ~ .,
+                             data = sonar_training_data)
>
> # View the forest results.
> print(output_fit)

Call:
randomForest(formula = X.1 ~ ., data = sonar_training_data)
Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 7

OOB estimate of error rate: 17.83%
Confusion matrix:
  -1  1 class.error
-1 56  9  0.1384615
 1 14 50  0.2187500
```

$$\begin{aligned}\text{Misclassification error rate} &= 1 - ((\text{True Positives} + \text{True Negatives}) / \text{Total observations}) \\ &= 1 - ((56 + 50) / (56 + 9 + 14 + 50)) \\ &= 1 - (106 / 129) \\ &= 1 - 0.8217054264 \\ &= \mathbf{0.1782945736}\end{aligned}$$

## 8) This question deals with sonar data

a) Use `knn()` for the k-nearest neighbor classifier for  $k=5$  and  $k=6$  to the last column of the sonar training data. Compute the misclassification error on the training data and also on the test data.

```
> setwd("E:\\MSIT-II\\Data Science Specialization\\Intro to ML Repo\\Data Mining\\DM Assignment4")
> sonar_train <- read.csv("sonar_train.csv")
> sonar_test <- read.csv("sonar_test.csv")
> x_train <- sonar_train[, 1:60]
> y_train <- sonar_train[, 61]
> x_test <- sonar_test[, 1:60]
> y_test <- sonar_test[, 61]
> names(sonar_train)
 [1] "x0.0258" "x0.0433" "x0.0547" "x0.0681" "x0.0784" "x0.125"  "x0.1296" "x0.1729"
 [9] "x0.2794" "x0.2954" "x0.2506" "x0.2601" "x0.2249" "x0.2115" "x0.127"  "x0.1193"
[17] "x0.1794" "x0.2185" "x0.1646" "x0.074"  "x0.0625" "x0.2381" "x0.4824" "x0.6372"
[25] "x0.7531" "x0.8959" "x0.9941" "x0.9957" "x0.9328" "x0.9344" "x0.8854" "x0.769"
[33] "x0.6865" "x0.639"  "x0.6378" "x0.6629" "x0.5983" "x0.4565" "x0.3129" "x0.4158"
[41] "x0.4325" "x0.4031" "x0.4201" "x0.4557" "x0.3955" "x0.2966" "x0.2095" "x0.1558"
[49] "x0.0884" "x0.0265" "x0.0121" "x0.0091" "x0.0062" "x0.0019" "x0.0045" "x0.0079"
[57] "x0.0031" "x0.0063" "x0.0048" "x0.005"  "x.1"

# argument 'c1' is missing, with no default
> knn_fit <- knn(x_train, x_train, k = 5, c1 = sonar_train$x.1)
> knn_fit
 [1] -1 -1 1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 1 -1 1 1 -1 -1 -1 1 1 -1
[27] 1 1 1 -1 1 -1 -1 -1 -1 -1 -1 -1 -1 1 -1 1 1 -1 1 -1 -1 -1 1
[53] -1 1 -1 -1 1 1 1 1 -1 -1 -1 -1 1 1 1 -1 -1 -1 -1 1 1 -1 -1
[79] 1 1 -1 1 -1 -1 1 -1 -1 -1 -1 -1 -1 1 -1 -1 1 -1 -1 -1 -1 -1
[105] 1 -1 -1 -1 1 1 -1 1 -1 -1 1 -1 -1 -1 -1 -1 1 -1 1 1 -1 1 -1
Levels: -1 1
> library(gmodels)
> CrossTable(x = sonar_train$x.1, y = knn_fit, prop.chisq = FALSE)
```



cell contents

		N
N / Row Total		
N / Col Total		
N / Table Total		

Total Observations in Table: 129

sonar_train\$x.1	knn_fit		Row Total
	-1	1	
-1	59	6	65
	0.908	0.092	0.504
	0.738	0.122	
	0.457	0.047	
1	21	43	64
	0.328	0.672	0.496
	0.263	0.878	
	0.163	0.333	
Column Total	80	49	129
	0.620	0.380	

```
> summary(knn_fit)
```

```
-1  1
```

```
80 49
```

```
> mis_error_train = 1-sum(knn_fit == y_train)/length(y_train)
```

```
> mis_error_train
```

```
[1] 0.2093023
```

```
> knn_fit <- knn(x_train, x_train, k = 6, cl = sonar_train$x.1)
```

```
> knn_fit
```

```
[1] -1 -1 1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 1 1 -1 -1 -1 1 1 -1
```

```
[27] 1 1 1 -1 1 -1 1 -1 -1 1 -1 -1 -1 -1 1 -1 1 -1 -1 1 -1 -1 -1 1
```

```
[53] 1 1 -1 -1 1 1 1 1 1 -1 -1 -1 -1 1 1 1 1 -1 -1 -1 1 1 -1 -1 -1
```

```
[79] 1 1 -1 1 -1 -1 1 -1 -1 -1 1 1 -1 -1 1 -1 1 -1 -1 -1 -1 -1 -1 -1
```

```
[105] -1 -1 -1 -1 1 1 -1 1 -1 -1 1 1 -1 -1 -1 -1 -1 1 1 -1 1 1 -1 1 -1
```

```
Levels: -1 1
```

```
> library(gmodels)
```

```
> CrossTable(x = sonar_train$x.1, y = knn_fit, prop.chisq = FALSE)
```

# Cell Contents

	N
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 129

sonar_train\$X.1	knn_fit		Row Total
	-1	1	
-1	57	8	65
	0.877	0.123	0.504
	0.750	0.151	
	0.442	0.062	
1	19	45	64
	0.297	0.703	0.496
	0.250	0.849	
	0.147	0.349	
Column Total	76	53	129
	0.589	0.411	

```
> mis_error_train_k6 = 1-sum(knn_fit == y_train)/length(y_train)
> mis_error_train_k6
[1] 0.2093023
```

```
> knn_fit_train_k5 <- knn(x_train, x_train, k = 5, cl = sonar_train$X.1)
> knn_fit_train_k5
[1] -1 -1 1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 1 -1 1 1 1 -1 -1 -1 1 1 -1
[27] 1 1 1 -1 1 -1 -1 -1 -1 -1 -1 -1 -1 1 -1 1 -1 -1 1 -1 -1 -1 1
[53] -1 1 -1 -1 1 1 1 1 -1 -1 -1 -1 1 1 1 1 -1 -1 -1 -1 1 1 -1 -1 -1
[79] 1 1 -1 1 -1 -1 1 -1 -1 -1 -1 1 -1 -1 1 -1 1 -1 -1 -1 -1 -1 -1
[105] 1 -1 -1 -1 1 1 -1 1 -1 -1 1 1 -1 -1 -1 -1 -1 1 -1 1 1 -1 1 -1
Levels: -1 1
> library(gmodels)
> CrossTable(x = sonar_train$X.1, y = knn_fit_train_k5, prop.chisq = FALSE)
```

# Cell Contents

	N
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 129

sonar_train\$x.1	knn_fit_train_k5		Row Total
	-1	1	
-1	59 0.908 0.738 0.457	6 0.092 0.122 0.047	65 0.504
1	21 0.328 0.263 0.163	43 0.672 0.878 0.333	64 0.496
Column Total	80 0.620	49 0.380	129

```
> misclassify_error_train_k5 = 1 - (sum(knn_fit_train_k5 == y_train) / length(y_train))
> misclassify_error_train_k5
[1] 0.2093023
```

```
> knn_fit_train_k6 <- knn(x_train, x_train, k = 6, cl = sonar_train$x.1)
> knn_fit_train_k6
 [1] -1 -1 1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 1 -1 1 1 1 -1 -1 -1 -1 1 -1
[27] 1 1 1 -1 1 -1 -1 -1 -1 1 -1 -1 -1 -1 1 -1 1 -1 -1 1 1 -1 1
[53] -1 1 -1 -1 1 1 1 1 -1 -1 -1 -1 1 1 1 1 1 -1 -1 -1 1 1 -1 -1 -1
[79] 1 1 -1 1 -1 -1 1 -1 -1 -1 -1 1 -1 -1 1 -1 1 -1 -1 -1 -1 -1 -1
[105] 1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 1 -1 -1 -1 -1 -1 1 1 1 1 1 -1
Levels: -1 1
> library(gmodels)
> CrossTable(x = sonar_train$x.1, y = knn_fit_train_k6, prop.chisq = FALSE)
```

# Cell Contents

	N
N / Row Total	
N / Col Total	
N / Table Total	

Total observations in Table: 129

sonar_train\$X.1	knn_fit_train_k6		Row Total
	-1	1	
-1	58 0.892 0.753 0.450	7 0.108 0.135 0.054	65 0.504
1	19 0.297 0.247 0.147	45 0.703 0.865 0.349	64 0.496
Column Total	77 0.597	52 0.403	129

```
> misclassify_error_train_k6 = 1 - (sum(knn_fit_train_k6 == y_train) / length(y_train))
> misclassify_error_train_k6
[1] 0.2015504
```

```
> knn_fit_test_k5 <- knn(x_test, x_test, k = 5, cl = sonar_test$X.1)
> knn_fit_test_k5
[1] -1 -1 1 -1 1 -1 1 -1 -1 -1 1 -1 1 -1 1 -1 -1 -1 -1 -1 -1 1 -1
[27] -1 1 1 1 -1 1 -1 -1 1 -1 -1 1 -1 1 -1 1 1 1 -1 -1 -1 1
[53] 1 -1 -1 -1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 -1 -1 -1 -1 1 1 -1 -1
Levels: -1 1
> library(gmodels)
> CrossTable(x = sonar_test$X.1, y = knn_fit_test_k5, prop.chisq = FALSE)
```

# Cell Contents

		N
N / Row Total		
N / Col Total		
N / Table Total		

Total Observations in Table: 77

sonar_test\$X.1	knn_fit_test_k5		Row Total
	-1	1	
-1	40 0.909 0.769 0.519	4 0.091 0.160 0.052	44 0.571
1	12 0.364 0.231 0.156	21 0.636 0.840 0.273	33 0.429
Column Total	52 0.675	25 0.325	77

```
> misclassify_error_test_k5 = 1 - (sum(knn_fit_test_k5 == y_test) / length(y_test))
> misclassify_error_test_k5
[1] 0.2077922
```

```
> knn_fit_test_k6 <- knn(x_test, x_test, k = 6, cl = sonar_test$X.1)
> knn_fit_test_k6
[1] -1 -1 1 -1 1 -1 1 -1 -1 -1 1 -1 1 -1 -1 -1 -1 -1 -1 1 -1
[27] 1 -1 1 1 -1 1 -1 1 -1 1 -1 1 1 1 -1 -1 -1 1 1 -1 -1 -1
[53] 1 -1 -1 -1 -1 -1 -1 1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 1 -1 -1
Levels: -1 1
> library(gmodels)
> CrossTable(x = sonar_test$X.1, y = knn_fit_test_k6, prop.chisq = FALSE)
```

Cell Contents	
	N
N / Row Total	
N / Col Total	
N / Table Total	

Total observations in Table: 77

sonar_test\$x.1	knn_fit_test_k6		Row Total
	-1	1	
-1	38 0.864 0.704 0.494	6 0.136 0.261 0.078	44 0.571
1	16 0.485 0.296 0.208	17 0.515 0.739 0.221	33 0.429
Column Total	54 0.701	23 0.299	77

```
> misclassify_error_test_k6 = 1 - (sum(knn_fit_test_k6 == y_test) / length(y_test))
> misclassify_error_test_k6
[1] 0.2857143
```

b) Repeat part a using the exact same R code a few times. Explain why both the training errors and the test errors often change for  $k=6$  but not for  $k=5$ . Hint: Read the help on the `knn` function if you do not know.

If there exists any tie in the values of  $K$ th nearest vector then all the candidates are included in vote.

For the value,  $k = 6$ , there are ties for 6<sup>th</sup> nearest vector.

Hence there are changes in the misclassification error rate as all the candidates will be included in the vote.